

TST_base_code

May 24, 2023

TST Representative Code This is a representative code for developing and training a full-scale encoder-decoder ‘Time-Series-Transformer’. Although the architecture is similar to the “Attention is all you need (Vaswani, et. al. 2017) paper, it has to be significantly modified to consider multi-dimensional float-type time-series sequences instead of one-dimensional textual inputs.

```
[ ]: import numpy as np
      from matplotlib import pyplot as plt
      import seaborn as sns
      import pandas as pd

      import torch
      import torch.nn as nn
      import random
```

How Attention Mechansim Works? Here is a tensorized description of how attention mechanism works along with the shapes of inputs. The output of the block is what you would get out of nn.Multiheadattention block from pytorch.

```
[ ]: import torch
      import torch.nn as nn

      class Attention(nn.Module):
          def __init__(self, d_model):
              super(Attention, self).__init__()
              self.query = nn.Linear(d_model, d_model)
              self.key = nn.Linear(d_model, d_model)
              self.value = nn.Linear(d_model, d_model)
              self.softmax = nn.Softmax(dim=-1)

          def forward(self, x):
              # Compute query, key, and value tensors
              q = self.query(x)
              k = self.key(x)
              v = self.value(x)

              # Compute attention scores
              attn_scores = torch.matmul(q, k.transpose(-2, -1))
              attn_scores = self.softmax(attn_scores)
```

```

    # Verify the sum of attention scores
    attn_sum = attn_scores.sum(dim=-1)
    print("Sum of attention scores:", attn_sum)

    # Apply attention weights to values
    attn_output = torch.matmul(attn_scores, v)

    return attn_output, attn_scores

# Example usage
batch_size = 1
seq_length = 12
d_model = 768

# Create a random input tensor
x = torch.randn(batch_size, seq_length, d_model)

# Create an instance of the Attention module
attention = Attention(d_model)

# Pass the input tensor through the attention mechanism
attn_output, attn_scores = attention(x)

# Print the shapes of the output tensors
print("Input shape:", x.shape)
print("Attention output shape:", attn_output.shape)
print("Attention scores shape:", attn_scores.shape)

```

```

Sum of attention scores: tensor([[1.0000, 1.0000, 1.0000, 1.0000, 1.0000,
1.0000, 1.0000, 1.0000, 1.0000,
1.0000, 1.0000, 1.0000]], grad_fn=<SumBackward1>)
Input shape: torch.Size([1, 12, 768])
Attention output shape: torch.Size([1, 12, 768])
Attention scores shape: torch.Size([1, 12, 12])

```

TST Hyperparameters

```

[ ]: window_size, prediction_horizon = 12, 6
    chunk_size = int(prediction_horizon/2)
    input_features, output_features = 6, 4

    num_encoder_blocks = 4
    num_decoder_blocks = 4
    nhead = 8
    dim_FFN = 512 # Dimensions of the FFN inside Encoder/Decoder Blocks
    dropout = 0.1

```

```
d_model = 768 # Vasani (2017)
batch_size = 32
```

Input Embedding Converting

$$X_{RAW} \rightarrow X_{EM}$$

$$[batch_size, W, input_features] \rightarrow [batch_size, W, d_model]$$

where d_{model} is the internal dimension of the transformer (lifted states), and W is the window size.

Reference: https://d2l.ai/chapter_attention-mechanisms-and-transformers/self-attention-and-positional-encoding.html#subsec-cnn-rnn-self-attention

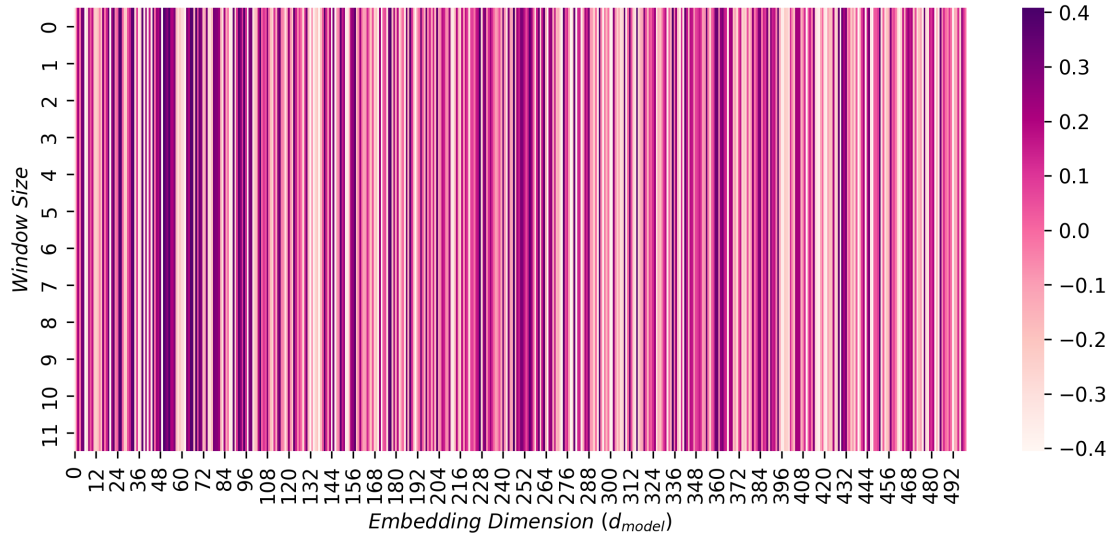
```
[ ]: class EmbeddingLayer(nn.Module):
    def __init__(self, input_size, d_model):
        super(EmbeddingLayer, self).__init__()
        self.fc = nn.Linear(input_size, d_model)

    def forward(self, x):
        x = self.fc(x)
        return x
```

```
[ ]: ## Example for encoding
sample = torch.zeros(batch_size, window_size, input_features)
encoder = EmbeddingLayer(input_features, d_model)
encoded_input = encoder(sample)
print('Raw Input [batch_size, window_size, input_features]: \n', sample.shape)
print('----- CONVERTED TO ----- ')
print('Encoded Input [batch_size, window_size, d_model]: \n', encoded_input.
      ↪shape)
```

```
Raw Input [batch_size, window_size, input_features]:
  torch.Size([32, 12, 6])
----- CONVERTED TO -----
Encoded Input [batch_size, window_size, d_model]:
  torch.Size([32, 12, 768])
```

```
[ ]: # Create a heatmap of the pos_encoded_input tensor
plt.figure(figsize=(10,4), dpi=300)
sns.heatmap(encoded_input[0,:,:500].detach().numpy(), cmap='RdPu')
# plt.xlabel('$Input~Features~(F_{in})$')
plt.xlabel('$Embedding~Dimension~(d_{model})$')
plt.ylabel('$Window~Size$')
# plt.title('Heat Map for Positional Encoding')
plt.show()
```



Positional Encoding Converting

$$X_{EM} \rightarrow X_{PE}$$

$$[batch_size, W, d_{model}] \rightarrow [batch_size, W, d_{model}]$$

```
[ ]: class PositionalEncoding(nn.Module):  #@save
    """Positional encoding."""
    def __init__(self, d_model, dropout, max_len=1000):
        super().__init__()
        self.dropout = nn.Dropout(dropout)
        # Create a long enough P
        self.P = torch.zeros((1, max_len, d_model))
        position = torch.arange(max_len, dtype=torch.float32).reshape(-1, 1)
        div_term = torch.pow(10000, torch.arange(0, d_model, 2, dtype=torch.
float32) / d_model)
        X = position/div_term

        # X = torch.arange(max_len, dtype=torch.float32).reshape(
        #     -1, 1) / torch.pow(10000, torch.arange(
        #         0, d_model, 2, dtype=torch.float32) / d_model)

        self.P[:, :, 0::2] = torch.sin(X)
        self.P[:, :, 1::2] = torch.cos(X)

    def forward(self, X):
        X = X + self.P[:, :X.shape[1], :].to(X.device)
        return self.dropout(X)
```

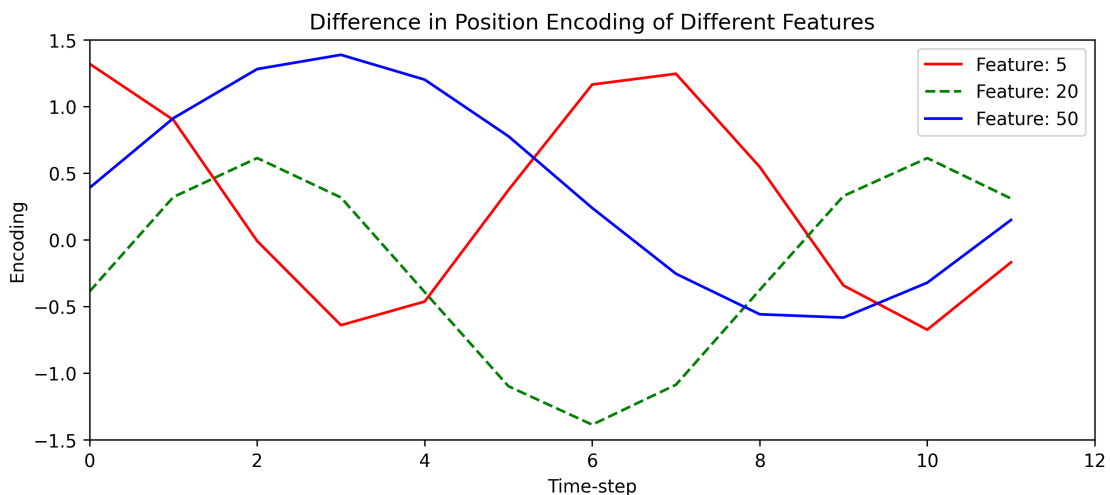
```
[ ]: pos_encoder = PositionalEncoding(d_model,0, max_len=window_size)
pos_encoded_input = pos_encoder(encoded_input)
print('Encoded Input [batch_size, window_size, d_model]: \n',encoded_input.
      ↪shape)
print('----- CONVERTED TO ----- ')
print('PE Input: [batch_size, window_size, d_model]:\n', pos_encoded_input.
      ↪shape)
```

```
Encoded Input [batch_size, window_size, d_model]:
  torch.Size([32, 12, 768])
----- CONVERTED TO -----
PE Input: [batch_size, window_size, d_model]:
  torch.Size([32, 12, 768])
```


Visual representation of positional encoding of different features

```
[ ]: # Define the colors and linestyles for each line
colors = ['red', 'green', 'blue']
linestyles = ['-', '--', '-']

# Plot each line with a different color and linestyle
plt.figure(figsize=(10,4), dpi=300)
for i, color, linestyle in zip([5, 20, 50], colors, linestyles):
    plt.plot(torch.arange(window_size), pos_encoded_input[0, :, i].detach(),
             ↪color=color, linestyle=linestyle, label=f'Feature: {i}')
plt.legend()
# Set the x-axis and y-axis limits
plt.xlim(0, window_size)
plt.ylim([-1.5,1.5])
plt.title('Difference in Position Encoding of Different Features')
plt.ylabel('Encoding')
plt.xlabel('Time-step')
plt.show()
```



```
[ ]: # Create a heatmap of the pos_encoded_input tensor
plt.figure(figsize=(10,4), dpi=300)
sns.heatmap(pos_encoded_input[0,:,:500].detach().numpy(), cmap='cool')
plt.xlabel('$Embedding~Dimension~(d_{model})$')
plt.ylabel('$Window~Size$')
plt.title('Heat Map for Positional Encoding')
plt.show()
```



Encoder Block

1. General Instructions:

Converting

$$X_{PE} \rightarrow X_{EN_i} \quad [for \ i = 1]$$

$$X_{EN_{i-1}} \rightarrow X_{EN_i} \quad [for \ i \neq 1]$$

Dimensions:

$$[batch_{size}, W, d_{model}] \rightarrow [batch_{size}, W, d_{model}]$$

where i is the i^{th} encoder block.

- 2. Layernorm:** LayerNorm (d_{model}) normalizes the input $[batch_{size}, W, d_{model}]$ in the d_{model} dimension to give out an output of the same dimensions $[batch_{size}, W, d_{model}]$.

3. Output of 'N' Encoder Blocks:

Any layer in the Transformer encoder does not change the shape of the input vectors. The positionally encoded input has the dimensions $[batch_size, W, d_{model}]$, and the output from all the encoded blocks has the dimension $[batch_size, W, d_{model}]$.

```
[ ]: class TransformerEncoderBlock(nn.Module):  #@save
      """The Transformer encoder block."""
      def __init__(self, d_model, nhead, dim_feedforward, dropout):
          super().__init__()
          self.self_attn = nn.MultiheadAttention(d_model, nhead,
          ↪ dropout=dropout, batch_first=True)
          self.feedforward = nn.Sequential(
              nn.Linear(d_model, dim_feedforward),
              nn.ReLU(),
              nn.Dropout(p=dropout),
              nn.Linear(dim_feedforward, d_model),
              nn.Dropout(p=dropout)
          )
          self.norm1 = nn.LayerNorm(d_model)
          self.norm2 = nn.LayerNorm(d_model)
          self.dropout1 = nn.Dropout(p=dropout)
          self.dropout2 = nn.Dropout(p=dropout)

      def forward(self, x):
          residual = x
          x = self.norm1(x)
          x, _ = self.self_attn(x, x, x)
          x = self.dropout1(x)
          x = residual + x


          residual = x
          x = self.norm2(x)
          x = self.feedforward(x)
          x = self.dropout2(x)
          x = residual + x
          return x
```

```
[ ]: encoder_block =
      ↪ TransformerEncoderBlock(d_model, nhead=nhead, dim_feedforward=dim_FFN, dropout=dropout)
      output = encoder_block(pos_encoded_input)
      print('Encoder Output [batch_size, window_size, d_model]:\n', output.shape)
```

```
Encoder Output [batch_size, window_size, d_model]:
  torch.Size([32, 12, 768])
```

Visualize the ‘transformation’ on the ‘positonally encoded input’ after an ‘Encoder Block’

```
[ ]: # Create a heatmap of the pos_encoded_input tensor
plt.figure(figsize=(10,4), dpi=300)
sns.heatmap(output[0,:,:500].detach().numpy(), cmap='plasma')
plt.xlabel('$Embedding-Dimension~(d_{model})$')
plt.ylabel('$Window-Size$')
plt.title('Heat Map after [1] Encoder Block')
plt.show()
```



Stacking Multiple Encoder Blocks

```
[ ]: class Multiple_Encoders(nn.Module):
    def __init__(self, input_size, output_size, d_model, nhead,
        ↪ dim_feedforward, num_encoder_blocks, w, h, dropout):
        super(Multiple_Encoders, self).__init__()

        self.w = w
        self.h = h

        self.encoder_blocks = nn.ModuleList([TransformerEncoderBlock(d_model,
        ↪ nhead, dim_feedforward, dropout=dropout) for _ in range(num_encoder_blocks)])

        self.encoder_linear = EmbeddingLayer(input_size, d_model)

        self.pos_encoder = PositionalEncoding(d_model, max_len = self.
        ↪ w, dropout=dropout)

    def forward(self, x):
        x = self.encoder_linear(x) # Pass through the encoder linear layer
```



```

x = self.pos_encoder(x) # Add positional encoding
for block in self.encoder_blocks:
    x = block(x) # Pass through the encoder TSTBlock modules
return x

```

Visualizing Output from Multiple Encoded Blocks

```


[ ]: multiple_encoder_blocks = Multiple_Encoders(input_size=input_features,output_size=output_features,
                                                d_model=d_model,nhead=nhead,dim_feedforward=dim_FFN,
                                                num_encoder_blocks=num_encoder_blocks,
                                                w=window_size,h=prediction_horizon,
                                                dropout=dropout)

encoder_output = multiple_encoder_blocks(sample)
print('Shape after Multiple Encoder Blocks:', encoder_output.shape)

# Create a heatmap of the pos_encoded_input tensor
plt.figure(figsize=(10,4), dpi=300)
sns.heatmap(encoder_output[0,:,:500].detach().numpy(), cmap='plasma')
plt.xlabel('$Embedding-Dimension~(d_{model})$')
plt.ylabel('$Window-Size$')
plt.title('Heat Map after [' + str(num_encoder_blocks) + '] Encoder Block')
plt.show()

```

Shape after Multiple Encoder Blocks: torch.Size([32, 12, 768])



Decoder Block

1. General Instructions:

Converting

$$\begin{aligned} X_{EN_i} &\rightarrow X_{DEC_1} \quad [for \ i = 1] \\ X_{DEC_{i-1}} &\rightarrow X_{DEC_i} \quad [for \ i \neq 1] \end{aligned}$$

Dimensions:

$$[batch_{size}, W, d_{model}] \rightarrow [batch_{size}, H + \lambda, d_{model}]$$

where W is the window size (used in the encoder side of things), and H is the prediction horizon (used in the decoder side of things).

2. Decoder Linear Layer and Positional Encoding

The decoder embedding layer takes in the input (y_{raw}) that has dimensions $[batch_size, H, F_out]$, and converts it to $[batch_size, H, d_{model}]$ represented by Y_{EM} . The positional encoder for the decoder side uses positional encoding with $max_len = (H + chunk_size)$ to convert it to Y_{PE} .

3. Output of ‘M’ Decoder Blocks:

The decoder takes in (y_{raw}) of size $[batch_size, H, F_{out}]$ and converts it to a decoder output of size $[batch_size, H, d_{model}]$.

4. Note:

The internal dimension of the transformer (d_{model}) should be the same for the encoder and decoder side. However, on encoder side the ‘seq_len’ is W , and on the decoder side it is H .

```
[ ]: class TransformerDecoderBlock(nn.Module):    #@save
      """The Transformer Decoder Block."""
      def __init__(self, d_model, nhead, dim_feedforward, dropout):
          super().__init__()
          self.self_attn = nn.MultiheadAttention(d_model, nhead,
          ↪ dropout=dropout, batch_first=True)
          self.cross_attn = nn.MultiheadAttention(d_model, nhead,
          ↪ dropout=dropout, batch_first=True)
          self.feedforward = nn.Sequential(
              nn.Linear(d_model, dim_feedforward),
              nn.ReLU(),
              nn.Dropout(p=dropout),
              nn.Linear(dim_feedforward, d_model),
              nn.Dropout(p=dropout)
          )
          self.norm1 = nn.LayerNorm(d_model)
          self.norm2 = nn.LayerNorm(d_model)
          self.norm3 = nn.LayerNorm(d_model)

          self.dropout1 = nn.Dropout(p=dropout)
```

```

self.dropout2 = nn.Dropout(p=dropout)
self.dropout3 = nn.Dropout(p=dropout)

def forward(self, x, memory):
    residual = x
    x = self.norm1(x)
    x, _ = self.self_attn(x, x, x)
    x = self.dropout1(x)
    x = residual + x

    residual = x
    x = self.norm2(x)
    x, cross_attn_weights = self.cross_attn(x, memory, memory)
    x = self.dropout2(x)
    x = residual + x

    residual = x
    x = self.norm3(x)
    x = self.feedforward(x)
    x = self.dropout3(x)
    x = residual + x
    return x, cross_attn_weights

```

```

[ ]: sample_decoder = torch.zeros(batch_size, prediction_horizon, d_model)
decoder_block=
↳TransformerDecoderBlock(d_model,nhead=nhead,dim_feedforward=dim_FFN,dropout=dropout)
# decoded_output = decoder_block(output[:,-prediction_horizon:,:
↳],sample_decoder)

decoded_output, cross_attn_weights = decoder_block(output[:,-prediction_horizon:
↳,:],output[:,-prediction_horizon:,:])
print('Decoder Output [batch_size, prediction_horizon, d_model]:
↳\n',decoded_output.shape)


# Create a heatmap
plt.figure(figsize=(10,4), dpi=300)
sns.heatmap(decoded_output[0,:,:500].detach().numpy(), cmap='viridis')
plt.xlabel('$Embedding-Dimension~(d_{model})$')
plt.ylabel('$Window-Size$')
plt.title('Heat Map after [1] Decoder Block')
plt.show()

```

```

Decoder Output [batch_size, prediction_horizon, d_model]:
torch.Size([32, 6, 768])

```



Stacking Multiple Decoder Blocks

```
[ ]: class Multiple_Decoders(nn.Module):
    def __init__(self, input_size, output_size, d_model, nhead,
    ↪dim_feedforward, num_decoder_blocks, w, h, chunk_size, dropout):
        super(Multiple_Decoders, self).__init__()

        self.w = w
        self.h = h

        self.decoder_blocks = nn.ModuleList([TransformerDecoderBlock(d_model,
    ↪nhead, dim_feedforward, dropout=dropout) for _ in range(num_decoder_blocks)])
        self.cross_attn_weights = [] # List to store cross-attention weights

        self.decoder_linear = EmbeddingLayer(output_size, d_model)

        self.pos_decoder = PositionalEncoding(d_model, max_len = self.h +
    ↪chunk_size, dropout=dropout)

    def forward(self, x, memory):
        x = self.decoder_linear(x) # Pass through the encoder linear layer
        x = self.pos_decoder(x) # Add positional encoding
        for block in self.decoder_blocks:
            x, cross_attn_weights = block(x, memory) # Pass through the
    ↪decoder TSTBlock modules
            self.cross_attn_weights.append(cross_attn_weights) # Save the
    ↪cross-attention weights
```

```
return x
```

Visualizing output from just the decoder:

First we need to take X_{RAW} and then run it through the encoder blocks, and then combine y_{RAW} for the cross-attention.

```
[ ]: sample_decoder = torch.zeros(batch_size, chunk_size+prediction_horizon,
    ↪output_features)

multiple_encoder_blocks =
    ↪Multiple_Encoders(input_size=input_features,output_size=output_features,
    ↪
    ↪d_model=d_model,nhead=nhead,dim_feedforward=dim_FFN
    ↪,num_encoder_blocks=num_encoder_blocks,
    ↪w>window_size,h=prediction_horizon,
    ↪dropout=dropout)


encoder_output = multiple_encoder_blocks(sample)
print('Shape after Multiple Encoder Blocks:', encoder_output.shape)

multiple_decoder_blocks = Multiple_Decoders(input_size=input_features,
    ↪output_size=output_features,
    ↪
    ↪d_model=d_model,nhead=nhead,dim_feedforward=dim_FFN,
    ↪
    ↪num_decoder_blocks=num_decoder_blocks,
    ↪w>window_size,h=prediction_horizon,
    ↪chunk_size=chunk_size,
    ↪dropout=dropout)

decoder_output = multiple_decoder_blocks(sample_decoder,encoder_output[:
    ↪,-(chunk_size+prediction_horizon):,:])
attention_weights = multiple_decoder_blocks.cross_attn_weights

# Create a heatmap of the pos_encoded_input tensor
plt.figure(figsize=(10,4), dpi=300)
sns.heatmap(decoder_output[0,:,:500].detach().numpy(), cmap='viridis')
plt.xlabel('$Embedding~Dimension~(d_{model})$')
plt.ylabel('$Window~Size$')
plt.title('Heat Map after [' + str(num_decoder_blocks) + '] Decoder Blocks')
plt.show()
```

Shape after Multiple Encoder Blocks: torch.Size([32, 12, 768])



Visualizing attention weights from multiple decoders

```
[ ]: batch_index = random.randint(0, batch_size - 1)
num_blocks = len(attention_weights)

# Create a single figure with subplots
fig, axes = plt.subplots(1, num_blocks, figsize=(num_blocks * 5, 10))

# Generate an array of alpha values
alpha_array = np.linspace(0.4, 1, num_blocks)

# Iterate over the attention weights and plot them
for i, weights in enumerate(attention_weights):
    # Select the attention weights for the specified batch
    weights = weights[batch_index, :, :]

    # Select the corresponding subplot
    ax = axes[i]

    # Plot the attention weights with the specified alpha value
    img = ax.imshow(weights.squeeze(0).detach().numpy(), cmap='hot',
interpolation='nearest')
    ax.set_title('Cross-Attention Weights (Decoder Block {})'.format(i + 1))
    ax.set_xlabel('Query')
    ax.set_ylabel('Key')

    # Set xlim and ylim to match weights.shape[-1]
    xlim = weights.shape[-1]
```

```


ylim = weights.shape[-1]
ax.set_xlim(1, xlim-1)
ax.set_ylim(1, ylim-1)

# Set xticks and yticks to display exact integers
ax.set_xticks(np.arange(0, xlim, 1))
ax.set_yticks(np.arange(0, ylim, 1))

# Add a color bar to the subplot
cbar = fig.colorbar(img, ax=ax, shrink=0.3)
cbar.ax.set_ylabel('Attention Score')

plt.tight_layout()
plt.show()

```



Output Module

1. General Instructions:

Converting decoder output to the model predictions (y_{tgt}) of dimensions $[batch_size, H + \lambda, F_{out}]$.

1. Entire Chain of Conversions for TSTs:

$$X_{RAW} \rightarrow X_{EM}$$

$$X_{EM} \rightarrow X_{PE}$$

$$X_{PE} \rightarrow X_{EN_i}$$

$$X_{EN_i} \rightarrow Y_{DEC_j}$$

$$Y_{DEC_j} \rightarrow Y_{tgt}$$

2. Dimensions:

The encoder embedding layer takes in input (X_{RAW}) that has dimensions $[batch_size, W, F_{in}]$, and converts it to $[batch_size, W, d_{model}]$ represented by X_{EM} . The positional encoder for the decoder side uses positional encoding with $max_len = W$ to convert it to X_{PE} . Then, X_{PE} is converted to X_{EN_i} through encoder blocks.

The decoder embedding layer takes in input (y_{raw}) that has dimensions $[batch_size, H + \lambda, F_{out}]$, and converts it to $[batch_size, H + \lambda, d_{model}]$ represented by

Y_{EM} The positional encoder for the decoder side uses positional encoding with $max_len = (H + chunk_size)$ to convert it to Y_{PE} . The the decoder blocks convert this to a tensor of dimensions $[batch_size, H + \lambda, d_{model}]$.

Finally, $[batch_size, H + \lambda, d_{model}]$ is converted to $[batch_size, H, F_{out}]$ (model predictions).

```
[ ]: output_FFN = EmbeddingLayer(d_model,output_features)
      final_output = output_FFN(decoder_output)
      final_output = final_output[:,-prediction_horizon:,:]

      print('Decoder Output: ', decoder_output.shape)
      print("Final Output: ", final_output.shape)
```

Decoder Output: torch.Size([32, 9, 768])

Final Output: torch.Size([32, 6, 4])

0.0.1 Combining all the modules for making a TST

```
[ ]: class TimeSeriesTransformer(nn.Module):
      def __init__(self, input_size, output_size, d_model=768, nhead=8,
      ↪dim_feedforward=512,
          num_encoder_blocks=6, num_decoder_blocks=6,
          dropout=0.1,
          window_size=window_size,
      ↪prediction_horizon=prediction_horizon, chunk_size = chunk_size):

          super(TimeSeriesTransformer, self).__init__()

          # Set the model parameters
          self.d_model = d_model
          self.nhead = nhead
          self.num_encoder_blocks = num_encoder_blocks
          self.num_decoder_blocks = num_decoder_blocks
          self.window_size = window_size
          self.prediction_horizon = prediction_horizon

          # Encoder blocks
          self.multiple_encoder_blocks =
      ↪Multiple_Encoders(input_size=input_features,output_size=output_features,
          ↪
      ↪d_model=d_model,nhead=nhead,dim_feedforward=dim_FFN
          ,num_encoder_blocks=num_encoder_blocks,
          w=window_size,h=prediction_horizon,
          dropout=dropout)

          # Decoder blocks
          self.multiple_decoder_blocks =
      ↪Multiple_Decoders(input_size=input_features, output_size=output_features,
```



```

        ↪ d_model=d_model,nhead=nhead,dim_feedforward=dim_FFN,
        ↪ num_decoder_blocks=num_decoder_blocks,
        ↪ chunk_size=chunk_size,
        ↪ w=window_size,h=prediction_horizon,
        ↪ dropout=dropout)

    # Output layer
    self.output_FFN = EmbeddingLayer(d_model,output_features)

    def forward(self, x_enc, x_dec):
        # Encode the input time series data
        encoder_output = self.multiple_encoder_blocks(x_enc)

        # Decode the input time series data
        decoder_output = self.multiple_decoder_blocks(x_dec, encoder_output[:,
        ↪ -(chunk_size+prediction_horizon):, :])

        # Pass the decoded output through the output layer
        final_output = self.output_FFN(decoder_output)

        # Return only the predicted values
        # final_output = final_output[:, -prediction_horizon:, :]

    return final_output

```

Testing if the Entire TST Block works

```

[ ]: model=
    ↪ TimeSeriesTransformer(input_size=input_features,output_size=output_features,d_model=d_model,
    ↪ nhead=nhead,dim_feedforward=dim_FFN,
    ↪ num_encoder_blocks=num_decoder_blocks,num_decoder_blocks=num_decoder_blocks,
    ↪ dropout=dropout>window_size=window_size,prediction_horizon=prediction_horizon,chunk_size=ch

input_sample = torch.zeros(batch_size, window_size, input_features)
output_sample = torch.zeros(batch_size, prediction_horizon, output_features)
decoder_input = torch.cat((input_sample[:, -chunk_size:, 1:-1], output_sample[:, :,
    ↪ chunk_size, :]), dim=1)

print('Encoder Input Dimensions: ', input_sample.shape)
print('Decoder Input Dimensions: ', decoder_input.shape)
print('TST Output Dimensions: ', output_sample.shape)

```

```
TST_output = model(input_sample,output_sample)
print(TST_output.shape)
```

```
Encoder Input Dimensions: torch.Size([32, 12, 6])
Decoder Input Dimensions: torch.Size([32, 6, 4])
TST Output Dimensions: torch.Size([32, 6, 4])
torch.Size([32, 6, 4])
```

0.0.2 Training Module for TST

Code snippet for training a TST from scratch given a train_loader, and val_loader dataset can be imported/generated.

```
[ ]: def train_model(model, train_loader, val_loader, loss_fn, optimizer, scheduler,
    ↪ num_epochs):
    train_losses = [] # List to store training losses for each epoch
    val_losses = [] # List to store validation losses for each epoch

    for epoch in range(num_epochs):
        # Training
        model.train() # Set the model to training mode
        train_loss = 0.0
        for i, (x_enc, x_dec, y) in enumerate(train_loader):
            optimizer.zero_grad() # Clear the gradients
            output = model(x_enc, x_dec) # Forward pass
            loss = loss_fn(output, y) # Compute the loss
            loss.backward() # Backward pass
            optimizer.step() # Update the model's parameters
            train_loss += loss.item() # Accumulate the training loss for the
    ↪ current batch
        scheduler.step() # Update the learning rate scheduler
        avg_train_loss = train_loss / len(train_loader) # Calculate the
    ↪ average training loss for the epoch
        train_losses.append(avg_train_loss) # Store the average training loss

        # Validation
        model.eval() # Set the model to evaluation mode
        val_loss = 0.0
        with torch.no_grad():
            for i, (x_enc, x_dec, y) in enumerate(val_loader):
                output = model(x_enc, x_dec) # Forward pass (no gradients
    ↪ needed)
                loss = loss_fn(output, y) # Compute the loss
                val_loss += loss.item() # Accumulate the validation loss for
    ↪ the current batch
            avg_val_loss = val_loss / len(val_loader) # Calculate the average
    ↪ validation loss for the epoch
```

```

        val_losses.append(avg_val_loss)  # Store the average validation loss

        # Print the epoch number and the corresponding training and validation
        ↪ losses
        print(f"Epoch {epoch+1}/{num_epochs}, train loss: {avg_train_loss:.4f},
        ↪ val loss: {avg_val_loss:.4f}")

        return train_losses, val_losses  # Return the training and validation
        ↪ losses for all epochs

```